

# APPLIED R FOR SOCIAL SCIENTISTS

---

J. Alexander Branham

Spring 2016

## THIS CLASS

---

- Assumption of *some* previous exposure to R
  - We're not explaining assignment, packages, calling functions, etc
- Will *not* be covering statistics
- Alex will start with some common data tasks: loading, variable creation, merging, etc
- Daniel will take the second part to talk about tables and visualization

## GETTING THE DATA

---

- `source()` runs a file through R
- This one checks if you have the data already and tries to download it if not
- The dataset we're using is the General Social Survey spanning 1972-2014

```
source("check-gss-and-maybe-download.R")
```

```
[1] "GSS file exists!"
```



package	function	file formats
foreign	<b>read.*</b>	dta, spss, etc
haven	<b>read_*</b>	dta, spss, etc
utils	<b>read.csv</b>	csv files
readr	<b>read_csv</b>	csv files

- We have stata data (\*.dta)
- Two packages in R: **foreign** and **haven**
- **foreign** is older but more stable
- **haven** is newer and can read Stata 13+ files
- **convert.factors=FALSE** ensures that R doesn't convert the values to the labels that stata uses

```
GSS <- foreign::read.dta("./data/GSS7214_R4.DTA",  
                        convert.factors = FALSE)
```



# DATA MANAGEMENT

---

## THE DATA: A TABLE

- Let's get a sense of the data we're working with
- Do you want more, less, or about the same spending? (education and social security)

```
table(GSS$nateduc, GSS$natsec, exclude = NULL)
```

	1	2	3	<NA>
1	8540	4335	699	7939
2	2062	2288	296	5005
3	353	457	210	1317
<NA>	11841	7408	1282	5567

- `cor` gives us (by default) the pearson's  $r$  between two variables
- Without setting `use`, R tries to use the whole data, some of which are missing and thus results in `NA`

```
cor(GSS$nateduc,  
    GSS$nat soc,  
    use="complete.obs")
```

```
[1] 0.1965368
```

- Let's make an indicator variable for whether a respondent is black or not
- Here's the **race** variable
- I also like to make sure that I'm not going to overwrite an existing variable

```
table(GSS$race)
```

1	2	3
48240	8312	3047

```
table(GSS$black)
```

```
< table of extent 0 >
```

- Using `ifelse` to create a variable conditional on other var's values

```
GSS$black <- with(GSS, ifelse(race == 2, TRUE, FALSE))  
table(GSS$black)
```

FALSE	TRUE
51287	8312

- Now let's check to see if that correlation is different for black people
- Note how ugly this looks!

```
cor(GSS$nateduc[GSS$black == TRUE],  
    GSS$nat soc[GSS$black == TRUE],  
    use="complete.obs")
```

```
[1] 0.1772446
```

- `dplyr` is an R package that makes data management *much* easier
- Different functions for data munging:
  - `filter()`, `select()`, `mutate()`
- It introduces the pipe operator `%>%` to the language
- Functions for merging data
  - `*_join`: full, inner, left, right
- `group_by`, which lets us perform operations on groups of the data

```
suppressPackageStartupMessages(library(dplyr))
```

## SUBSETTING THE DPLYR WAY

- The pipe (`%>%`) “pipes” the output of the last thing into the first argument of the next thing
- `summarize` (or `summarise`) from `dplyr` returns a `data.frame`

```
with(filter(GSS, black == TRUE),  
      cor(nateduc, natsoc,  
          use = "complete.obs"))
```

```
[1] 0.1772446
```

```
GSS %>%  
  filter(black == TRUE) %>%  
  summarize(mycor =  
            cor(nateduc, natsoc,  
                use = "complete.obs"))
```

```
mycor  
1 0.1772446
```



## DROPPING OBSERVATIONS

- 1972 doesn't have any observations we're interested in (our spending variables weren't asked), so let's drop it
- Again, we can use `filter`, but this time we assign the result back to `GSS`:

```
GSS <- GSS %>%  
  filter(year != 1972)
```

## FACTORS

- Variables with categories can be represented as factors in R
- If you want R to think they're ordered, you can use `ordered = TRUE` as an argument

```
table(GSS$sex)
```

```
GSS <- GSS %>%
```

```
  mutate(sex = factor(sex,  
                        levels = c(1,2),  
                        labels = c("M","F")))
```

```
      1      2  
25479 32507
```

```
table(GSS$sex)
```

```
      M      F  
25479 32507
```

- `dplyr` provides `group_by`
- Lets us perform operations to grouped data

```
thecors <- GSS %>%  
  group_by(sex, black) %>%  
  summarize(thecor = cor(nateduc, natsoc,  
                        use = "complete.obs"),  
            n = n())
```

```
print(thecors)
```

Source: local data frame [4 x 4]

Groups: sex [?]

	sex	black	thecor	n
	(fctr)	(lg1)	(dbl)	(int)
1	M	FALSE	0.1918454	22446
2	M	TRUE	0.1674413	3033
3	F	FALSE	0.1786193	27489
4	F	TRUE	0.1820090	5018

- Maybe we're interested in preferences by year?

```
gss_yearly <- GSS %>%  
  group_by(year) %>%  
  summarize(educ = mean(nateduc,  
                        na.rm=TRUE),  
            soc = mean(natsoc,  
                      na.rm=TRUE))
```

```
head(gss_yearly)
```

```
Source: local data frame [6 x 3]
```

	year (int)	educ (dbl)	soc (dbl)
1	1973	1.582287	NaN
2	1974	1.562059	NaN
3	1975	1.604930	NaN
4	1976	1.579020	NaN
5	1977	1.605854	NaN
6	1978	1.576766	NaN

- Means are nice, but there are other ways to summarize data
- What if we want to look at the proportion of people who support more spending minus the proportion who support less?

```
netsupport <- function(thedata){  
  prop_more <- mean(thedata == 1, na.rm = TRUE)  
  prop_less <- mean(thedata == 3, na.rm = TRUE)  
  prop_more - prop_less  
}
```

```
GSS %>%  
  group_by(year) %>%  
  summarize(support_educ=netsupport(nateduc),  
            support_soc=netsupport(natsoc))
```

Source: local data frame [29 x 3]

	year (int)	support_educ (dbl)	support_soc (dbl)
1	1973	0.4177127	NA
2	1974	0.4379408	NA
3	1975	0.3950704	NA
4	1976	0.4209800	NA
5	1977	0.3941457	NA



- The **ggplot2** package provides the **economics** data.frame that has US economic data starting in July 1967
- `?economics` gives more info

```
library(ggplot2)
```

```
head(economics)
```

```
Source: local data frame [6 x 6]
```

	date	pce	pop	psavert	uempmed	unemploy
	(date)	(dbl)	(int)	(dbl)	(dbl)	(int)
1	1967-07-01	507.4	198712	12.5	4.5	2944
2	1967-08-01	510.5	198911	12.5	4.7	2945
3	1967-09-01	516.3	199113	11.7	4.6	2958
4	1967-10-01	512.9	199311	12.5	4.9	3143
5	1967-11-01	518.1	199498	12.5	4.7	3066

- Let's make an unemployment rate by `unemploy/pop`

```
economics <- economics %>%  
  mutate(unemp_rate = unemploy / pop)
```

- The `economics` data is monthly and our GSS data is yearly, so we need to aggregate

```
economics_yearly <- economics %>%  
  mutate(year = format(date, "%Y")) %>%  
  group_by(year) %>%  
  summarize(unemp = mean(unemp_rate))
```

- Let's see what our data looks like now!

```
head(economics_yearly)
```

```
Source: local data frame [6 x 2]
```

	year (chr)	unemp (dbl)
1	1967	0.01512179
2	1968	0.01394202
3	1969	0.01396464
4	1970	0.02012547
5	1971	0.02418970
6	1972	0.02323808

- Now we have two `data.frame` objects — `gss_yearly` and `economics_yearly` — that we want to join together
- `dplyr` provides a really easy way of doing this
- The jargon comes from SQL, a programming language used to store data
- What you probably call a “merge” `dplyr` calls a “join”
- `*_join` where `*` is either `full`, `inner`, `left`, or `right`
- We’ll use `left_join` since the economics data contains years that aren’t in the GSS

```
gss_yearly <- left_join(gss_yearly,  
                        economics_yearly,  
                        by = "year")
```

Error: cannot join on columns 'year' x 'year': Can't join on 'year'

- Error: cannot join on columns 'year' x 'year': Can't join on 'year' x 'year' because of incompatible types (character / integer)

- The error on the last slide indicates that the **year** variable in the two datasets is different
- Let's verify that:

```
class(gss_yearly$year)
```

```
[1] "integer"
```

```
class(economics_yearly$year)
```

```
[1] "character"
```



- Solution: change `economics_yearly$year` to an integer

```
economics_yearly$year <- as.integer(economics_yearly$year)
```

```
gss_yearly <- left_join(gss_yearly,  
                        economics_yearly,  
                        by="year")
```

```
head(gss_yearly)
```

```
Source: local data frame [6 x 4]
```

	year (int)	educ (dbl)	soc (dbl)	unemp (dbl)
1	1973	1.582287	NaN	0.02057710
2	1974	1.562059	NaN	0.02418823
3	1975	1.604930	NaN	0.03677624
4	1976	1.579020	NaN	0.03393594
5	1977	1.605854	NaN	0.03164388
6	1978	1.576766	NaN	0.02780555

- Maybe you want to save this new data so you don't have to re-run the merging whenever you want to

package	function	result
readr	<b>write_csv</b>	csv file
utils	<b>write.csv</b>	csv file
base	<b>save</b>	Rdata file
xlsx	<b>write.xlsx</b>	excel file

- R can also write to stata/SPSS/SAS files through **foreign** or **haven**

- Let's save a csv file
- If the **data/** subfolder doesn't exist, this will produce an error
- The script that we ran at the beginning created this if it didn't already exist

```
readr::write_csv(gss_yearly, "data/gss-yearly-data.csv")
```